What is a Multi-Stage Build in Docker?

Multi-stage build allows you to use multiple FROM statements in a single Dockerfile to:

- Build the app in one stage
- Copy only what's needed to a smaller final image

? Why Do We Need It?

✓ Main Goals:

𝚱 Benefit	
♦ Smaller Images	Only copy what's needed into final image
ெ More Secure	No dev tools or secrets in production image
& Cleaner CI/CD	Separate build & runtime environment
Better Layer Caching	Speeds up builds
Environment Separation	One image builds everything!

Real-World Analogy

Imagine:

- **E** Stage 1 = Construction site (messy, heavy tools)
- Stage 2 = Finished house (clean, cozy)

You **build** in the messy environment, but **only move the furniture** into the clean house. \mathcal{A}

Multi-Stage Build Syntax

```
COPY --from=builder /app/package.json ./
RUN npm ci --omit=dev

# Set env vars, port and run
ENV PORT=3000
EXPOSE 3000
CMD ["node", "dist/index.js"]
```

Key Concepts Explained

Keyword	Meaning
AS builder	Give a name to this stage
from=builder	Copy files from previous stage
npm ciomit=dev	Install only production deps
COPY	Used only in build stage to avoid code bloat in final image

Before vs After: Image Size

Approach	Image Size	Contents
Traditional Single Build	~900MB	Full source code + dev dependencies
Multi-Stage Build	~200MB	Just built app + runtime dependencies

Real Project Example: React App

```
# Step 1: Build React App
FROM node:20-alpine AS builder
WORKDIR /app
COPY package*.json ./
RUN npm install
COPY . .
RUN npm run build

# Step 2: Serve using NGINX
FROM nginx:alpine
COPY --from=builder /app/build /usr/share/nginx/html
EXPOSE 80
CMD ["nginx", "-g", "daemon off;"]
```

☑ This builds the app with Node.js, and serves the static files via NGINX (no Node.js in final image!)

& Common Multi-Stage Use Cases

Use Case	Description
Frontend builds	Use node + nginx combo
% Backend builds	Build with TS/Go/Rust, then copy binaries only
Testing stage	Add test/linting in one stage, skip in final
CI/CD pipelines	Clean, reproducible builds across stages

Pro Tips & Best Practices

♀ Tip	✓ Recommendation
Useomit=dev	Strip dev-only packages in final stage
Use .dockerignore	Exclude node_modules, .git, tests/, etc
Use labels	Add metadata like version, author, etc
Don't copy everything	Use exact COPY paths for size control
Use named stages	Easier to copy from (from=builder)
Keep final image minimal	Just enough to run your app (no tools!)

Combine with Docker Compose

You can define multi-stage builds in your Dockerfile and just run:

```
docker-compose build
docker-compose up
```

Your services will use the optimized final image automatically a

Example Multi-Stage for TypeScript API

```
# Stage 1: Compile TS
FROM node:20-alpine AS builder
WORKDIR /app
COPY . .
RUN npm install
RUN npm run build

# Stage 2: Run with only JS output
FROM node:20-alpine
WORKDIR /app
COPY --from=builder /app/dist ./dist
COPY --from=builder /app/package*.json ./
```

```
RUN npm ci --omit=dev
CMD ["node", "dist/server.js"]
```

- ✓ Summary: When to Use Multi-Stage Builds?
- Always use if:
 - You're using build tools like tsc, webpack, vite
 - You want minimal production images
 - You want to separate testing/staging/building
 - You want faster CI builds & smaller attack surface

Final TL;DR Cheatsheet

Stage	Purpose	Base Image	Output
Stage 1 (builder)	Build, compile, test	node, golang, rust, etc.	/dist,/build,etc.
Stage 2 (prod)	Serve/run app only	node:alpine, nginx, etc.	Final slim image

Full Dockerfile (Context Recap)

```
# Stage 1: Build Stuff
FROM base as builder

WORKDIR /home/build

COPY package*.json .

COPY tsconfig.json .

RUN npm install

COPY src/ src/

RUN npm run build

# Stage 2: Runner
FROM base as runner

WORKDIR /home/app

COPY --from=builder /home/build/dist dist/
COPY --from=builder /home/build/package*.json .

RUN npm install --omit=dev
```

```
RUN addgroup --system --gid 1001 nodejs
RUN adduser --system --uid 1001 nodejs

USER nodejs

EXPOSE 8000
ENV PORT=8000

CMD [ "npm", "start" ]
```

& A2Z Breakdown of Each Section

FROM node:20-alpine3.19 as base

What it does:

- Starts from a minimal Node.js 20 Alpine image
- Alpine is lightweight (~5MB), good for small, fast images
- as base names this stage for reuse
- Think of base like a shared template that both stages use.

≪ Stage 1: Builder

```
FROM base as builder
WORKDIR /home/build
```

What happens here:

- We switch to a new build stage, using base image
- WORKDIR /home/build sets a directory for our build process

```
COPY package*.json .
COPY tsconfig.json .
RUN npm install
```

(f) Install dependencies:

- package*.json copied to install dependencies
- tsconfig.json is required for TypeScript compilation
- npm install installs all dependencies (dev + prod)

```
COPY src/ src/
RUN npm run build
```

% Build your app:

- Copies your app's TypeScript code
- npm run build compiles TS into JS → typically inside /dist

✓ End Result of Stage 1:

A folder /home/build/dist with compiled production-ready JS output.

Stage 2: Runner

```
FROM base as runner WORKDIR /home/app
```


- We now create a fresh container just for **running** the app.
- WORKDIR /home/app is where your app will run from.

```
COPY --from=builder /home/build/dist dist/
COPY --from=builder /home/build/package*.json .
```

Copy built artifacts only:

- Only copy the dist/ folder and package files (no source, no tsconfig)
- Ensures the final image is slim & clean

```
RUN npm install --omit=dev
```

Production-only install:

- Installs only prod dependencies (no dev tools like eslint, tsc, etc.)
- Keeps final image light and secure

Add Secure Non-Root User

```
RUN addgroup --system --gid 1001 nodejs
RUN adduser --system --uid 1001 nodejs
USER nodejs
```

⊕ Why?

- Running as root is dangerous in containers X
- We create a user nodejs with limited permissions for safety
- UID/GID 1001 is just an arbitrary non-root system user

Port & Env Setup

```
EXPOSE 8000
ENV PORT=8000
```

- EXPOSE 8000: Documents that the app uses port 8000
- ENV PORT=8000: Sets the default port for app to use internally

You still need to use -p to map it to host: docker run -p 8000:8000 <image>

Start the App

```
CMD [ "npm", "start" ]
```

Default entrypoint when container runs

• This triggers your "start" script from package.json:

```
"start": "node dist/index.js"
```

✓ Summary Table

♦ Section	Q Purpose
FROM base	Reuse image to reduce duplication
builder	Compiles TypeScript into JS
runner	Runs a minimal production image
npm install in builder	Installs full deps for building

Purpose

	· 1
npm installomit=dev in runner	Installs only what's needed to run
COPYfrom=builder	Efficient file copy without rebuild
USER nodejs	Enhances container security

Resulting Benefits

⋄ Section

	✓ Achieved
Small Image	Only runtime code in final image
Secure	✓ Non-root user, no dev tools
Faster Builds	✓ Reuses build layers
Clean Code Separation	✓ No TypeScript or build files inside final container
Portable	☑ Can run on any platform with Node 20

Bonus Tip: View Image Sizes

docker images

Compare the multi-stage image (~100MB) vs a single-stage image (~400–600MB)

Final Thoughts

This approach follows **Docker best practices**:

- Multi-stage
- Production-ready
- Secure by default
- Reproducible builds